

Project 1: regression 1

Urvan Christen, Amandine Goffeney, Joseph Vermeil, Lucile Vigu?f©


March 20, 2019

Introduction

Mercury is a metal present in the environment whose harmful effects on human health are well assessed (Park and Zheng 2012). In a study led in 2000 (Al-Majed and Preston 2000), scientists collected data on total mercury and methyl mercury levels in the hair of 100 fishermen of Kuwait, aged 16 to 58 years, comparing them to those of a control population of 35 non-fishermen, aged 26 to 35 years. The aim of this report is to analyse the factors influencing the levels of mercury in both populations. For the sake of simplicity, we will only focus on total Hg, leaving out methyl mercury, as both variables are strongly correlated.

The dataset gathers information about six numerical variables (age, height, weight, number of fish meals per week and residence time in Kuwait) and two categorical ones (being a fisherman or not, fish consumption habits). There is no additional information about gender as all the participants in the study are males.

A first insight at the data shows that the fishermen population exhibits higher levels of mercury in their hair. The significance of the difference between the means of both distributions is assessed by a *Welch Two Sample t-test* with the alternative hypothesis that the fishermen population has a average greater level of mercury than the control population (p-value = 7.473e-05).




Exploratory analysis (Amandine)

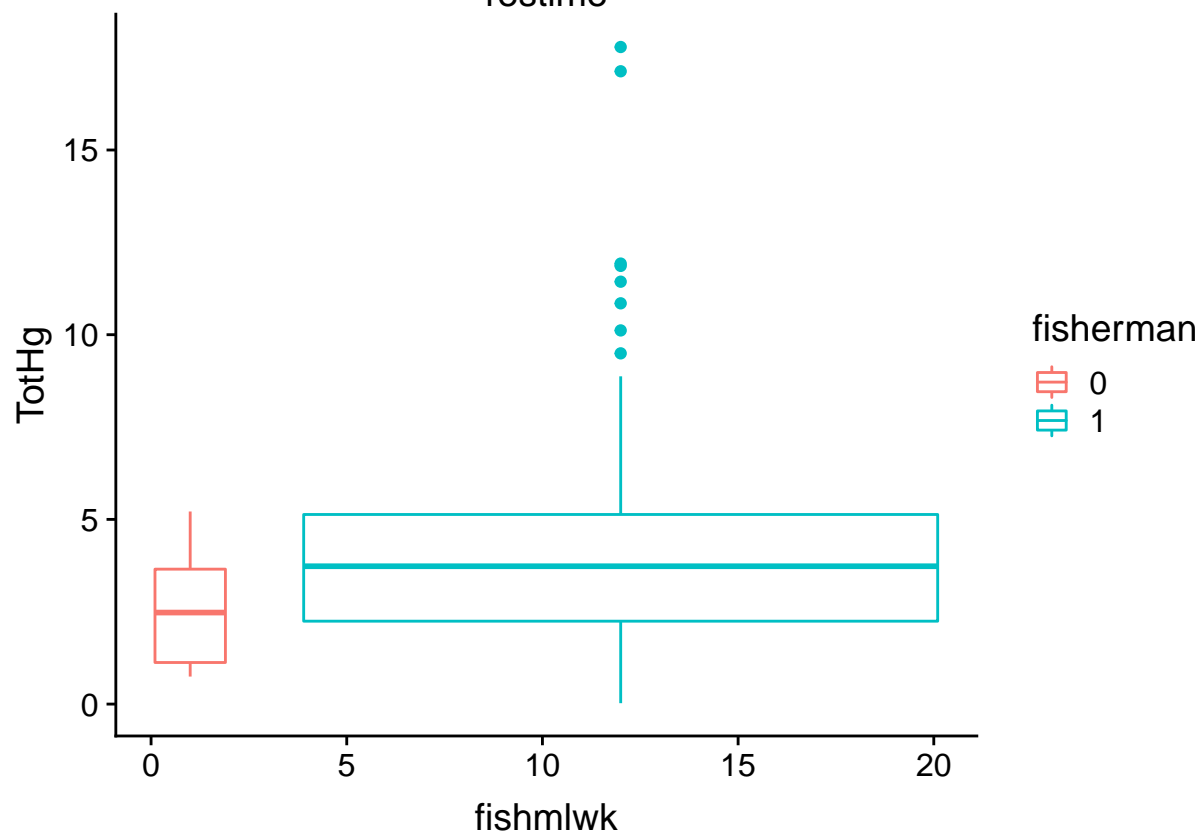

Overview of the data

Before fitting a model, let's have a look at the data. The following table shows the 8 possible values of `fishmlwk` according to the 2 possible values of `fisherman`.

```
##
##      0  1  2  3  4  7 14 21
##    0 10 14 11  0  0  0  0
##    1  0  0  0  2 12 70  5 11
```

First, let's notice that for some values, we have a very few people, for instance only 2 people eat fish 3 times a week. We see also that the data is really unbalanced, because the fishermen and non fishermen are separated on the variable “`fishmlwk`”: the non fishermen don't eat fish more than 2 times a week, whereas the fishermen don't eat fish less than 3 times a week. Besides, we face the same unbalanced pattern for the variables `age` and `restime`.





Multicollinearity

To check whether we have multicollinear variables, we use the VIF function. If all the variables have a result below 5, we keep them.

```
##              GVIF Df GVIF^(1/(2*Df))
## fisherman 2.176756 1      1.475383
## age       1.609314 1      1.268587
## restime   1.609630 1      1.268712
## height    1.138614 1      1.067059
## weight    1.220641 1      1.104826
## fishmlwk  1.754183 1      1.324456
## fishpart  1.745741 3      1.097312
```

Indeed, we observe that all the variables have a VIF below 2, so we don't eliminate any variable, yet.

Model selection (Joseph)

Now we are going to use and compare the different methods of model selection to select the more relevant parameters to explain the TotHg variations within the population.

Stepwise selection

The first attempt is a stepwise selection based on a formula with all the parameters except for "fisherman". Indeed, the preliminary exploration has shown that the "fisherman" variable have considerable impact on many of the other variables. Therefore to understand what is the reason behind the TotHg difference between the 2 groups, it seems natural to start with a model without the "fisherman" variable.

```
## Start:  AIC=260.18
## TotHg ~ age + restime + height + weight + fishmlwk + fishpart
##
##              Df Sum of Sq    RSS    AIC
## - height      1      4.793 816.53 258.97
## - age          1      8.814 820.55 259.63
## - fishpart     3     33.660 845.39 259.66
## - restime      1     10.460 822.19 259.90
## <none>                    811.73 260.18
## - fishmlwk     1     46.822 858.55 265.75
## - weight       1    115.001 926.73 276.06
##
## Step:  AIC=258.97
## TotHg ~ age + restime + weight + fishmlwk + fishpart
```

```

##
##           Df Sum of Sq    RSS    AIC
## - age      1      9.289 825.81 258.50
## - fishpart  3     36.427 852.95 258.86
## <none>                        816.53 258.97
## - restime   1     12.200 828.73 258.97
## + height    1      4.793 811.73 260.18
## - fishmlwk  1     45.206 861.73 264.25
## - weight    1    139.057 955.58 278.20
##
## Step:  AIC=258.5
## TotHg ~ restime + weight + fishmlwk + fishpart
##
##           Df Sum of Sq    RSS    AIC
## - restime   1      4.799 830.61 257.28
## <none>                        825.81 258.50
## + age       1      9.289 816.53 258.97
## - fishpart  3     40.857 866.67 259.02
## + height    1      5.267 820.55 259.63
## - fishmlwk  1     52.912 878.73 264.88
## - weight    1    135.101 960.92 276.95
##
## Step:  AIC=257.28
## TotHg ~ weight + fishmlwk + fishpart
##
##           Df Sum of Sq    RSS    AIC
## - fishpart  3     37.460 868.07 257.24
## <none>                        830.61 257.28
## + height    1      6.330 824.28 258.25
## + restime   1      4.799 825.81 258.50
## + age       1      1.888 828.73 258.97
## - fishmlwk  1     49.571 880.19 263.11
## - weight    1    133.220 963.83 275.36
##
## Step:  AIC=257.24
## TotHg ~ weight + fishmlwk
##
##           Df Sum of Sq    RSS    AIC
## <none>                        868.07 257.24
## + fishpart  3     37.460 830.61 257.28
## + height    1      9.170 858.90 257.80
## + age       1      5.744 862.33 258.34
## + restime   1      1.402 866.67 259.02
## - fishmlwk  1     95.208 963.28 269.29
## - weight    1    182.853 1050.93 281.04

```

```
##
## Call:
## lm(formula = TotHg ~ weight + fishmlwk, data = dataset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8344 -1.3096 -0.2953  0.6279 11.8572
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -10.07682    2.44481  -4.122 6.60e-05 ***
## weight       0.17518    0.03322   5.273 5.34e-07 ***
## fishmlwk     0.15884    0.04175   3.805 0.000216 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.564 on 132 degrees of freedom
## Multiple R-squared:  0.2498, Adjusted R-squared:  0.2384
## F-statistic: 21.97 on 2 and 132 DF,  p-value: 5.787e-09
```

With stepwise selection, the selected model is : TotHg ~ weight + fishmlwk. The intercept and the two coefficients are very significant. Moreover the signs of the coefficients are not absurd: while it is not really intuitive that the weight coefficient should be positive or negative, the coefficient of fishmlwk had to be positive, and it's the case here.

Now to determine possible differences between the fishermen and non-fishermen, a model based on the interactions between the “fisherman” variable and all the others is proposed. This could show if a certain variable is very relevant concerning fishermen but less when it comes to non-fishermen, for instance.

```
## Start:  AIC=269.38
## TotHg ~ fisherman:(age + restime + height + weight + fishmlwk +
##      fishpart)
##
##              Df Sum of Sq    RSS    AIC
## - fisherman:fishpart  5    19.102 802.53 262.64
## - fisherman:height    2     5.803 789.23 266.38
## - fisherman:age       2     7.850 791.28 266.73
## - fisherman:restime   2    14.368 797.80 267.84
## - fisherman:fishmlwk  2    23.242 806.67 269.33
## <none>                    783.43 269.38
## - fisherman:weight    2   123.792 907.22 285.19
##
## Step:  AIC=262.64
## TotHg ~ fisherman:age + fisherman:restime + fisherman:height +
##      fisherman:weight + fisherman:fishmlwk
```

```
##
##              Df Sum of Sq    RSS    AIC
## - fisherman:age      2    10.831 813.36 260.45
## - fisherman:restime  2    11.159 813.69 260.50
## - fisherman:height   2    12.973 815.50 260.80
## <none>                      802.53 262.64
## - fisherman:fishmlwk  2    44.977 847.51 266.00
## + fisherman:fishpart  5    19.102 783.43 269.38
## - fisherman:weight    2   140.813 943.34 280.46
##
## Step:  AIC=260.45
## TotHg ~ fisherman:restime + fisherman:height + fisherman:weight +
##      fisherman:fishmlwk
##
##              Df Sum of Sq    RSS    AIC
## - fisherman:restime  2     3.472 816.83 257.02
## - fisherman:height   2    16.323 829.69 259.13
## <none>                      813.36 260.45
## + fisherman:age      2    10.831 802.53 262.64
## - fisherman:fishmlwk  2    51.711 865.07 264.77
## + fisherman:fishpart  5    22.083 791.28 266.73
## - fisherman:weight    2   138.526 951.89 277.68
##
## Step:  AIC=257.02
## TotHg ~ fisherman:height + fisherman:weight + fisherman:fishmlwk
##
##              Df Sum of Sq    RSS    AIC
## - fisherman:height   2    16.652 833.49 255.75
## <none>                      816.83 257.02
## + fisherman:restime  2     3.472 813.36 260.45
## + fisherman:age      2     3.144 813.69 260.50
## - fisherman:fishmlwk  2    51.841 868.68 261.33
## + fisherman:fishpart  5    18.081 798.75 264.00
## - fisherman:weight    2   135.297 952.13 273.71
##
## Step:  AIC=255.75
## TotHg ~ fisherman:weight + fisherman:fishmlwk
##
##              Df Sum of Sq    RSS    AIC
## <none>                      833.49 255.75
## + fisherman:height   2    16.652 816.83 257.02
## - fisherman:fishmlwk  2    43.292 876.78 258.58
## + fisherman:age      2     6.170 827.32 258.74
## + fisherman:restime  2     3.800 829.69 259.13
## + fisherman:fishpart  5    26.422 807.06 261.40
```

```
## - fisherman:weight      2    204.147 1037.63 281.32
##
## Call:
## lm(formula = TotHg ~ fisherman:weight + fisherman:fishmlwk, data = dataset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.0705 -1.1391 -0.2026  0.6525 11.4266
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -9.26610     2.48392  -3.730 0.000284 ***
## fisherman0:weight  0.14464     0.03671   3.939 0.000133 ***
## fisherman1:weight  0.17361     0.03419   5.078 1.29e-06 ***
## fisherman0:fishmlwk 1.09393     0.60324   1.813 0.072076 .
## fisherman1:fishmlwk 0.09691     0.05274   1.837 0.068420 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.532 on 130 degrees of freedom
## Multiple R-squared:  0.2797, Adjusted R-squared:  0.2575
## F-statistic: 12.62 on 4 and 130 DF,  p-value: 1.053e-08
```

The best model is now : $\text{TotHg} \sim \text{fisherman:weight} + \text{fisherman:fishmlwk}$, ie the interactions between fisherman and the variables of the previous model. While the coefficients for fisherman0:weight and fisherman1:weight stays very significant, fisherman0:fishmlwk and fisherman1:fishmlwk are barely significant (but their p-value is very close to 0.05) ; this comes from the fact that within each of the subpopulation {fishermen} and {non-fishermen} the range in which fishmlwk varies is very reduced compared to the whole population (see the Exploratory analysis). Moreover, when looking at the values of the coefficients, it appears that for the non-fishermen, the number of fish meal per week has a very preponderant effect, while among fishermen, the contribution of weight and fish meal per week are almost even.

In order to check the dependency of the selected model to the selection technique, we have tried backward and forward selection and obtained similar models.

Results and discussion (Urvan)

Fit model (* fishermen) (I picked the model given by stepwise selection)


```
##
## Call:
## lm(formula = hg.form.selected, data = dataset)
##
```

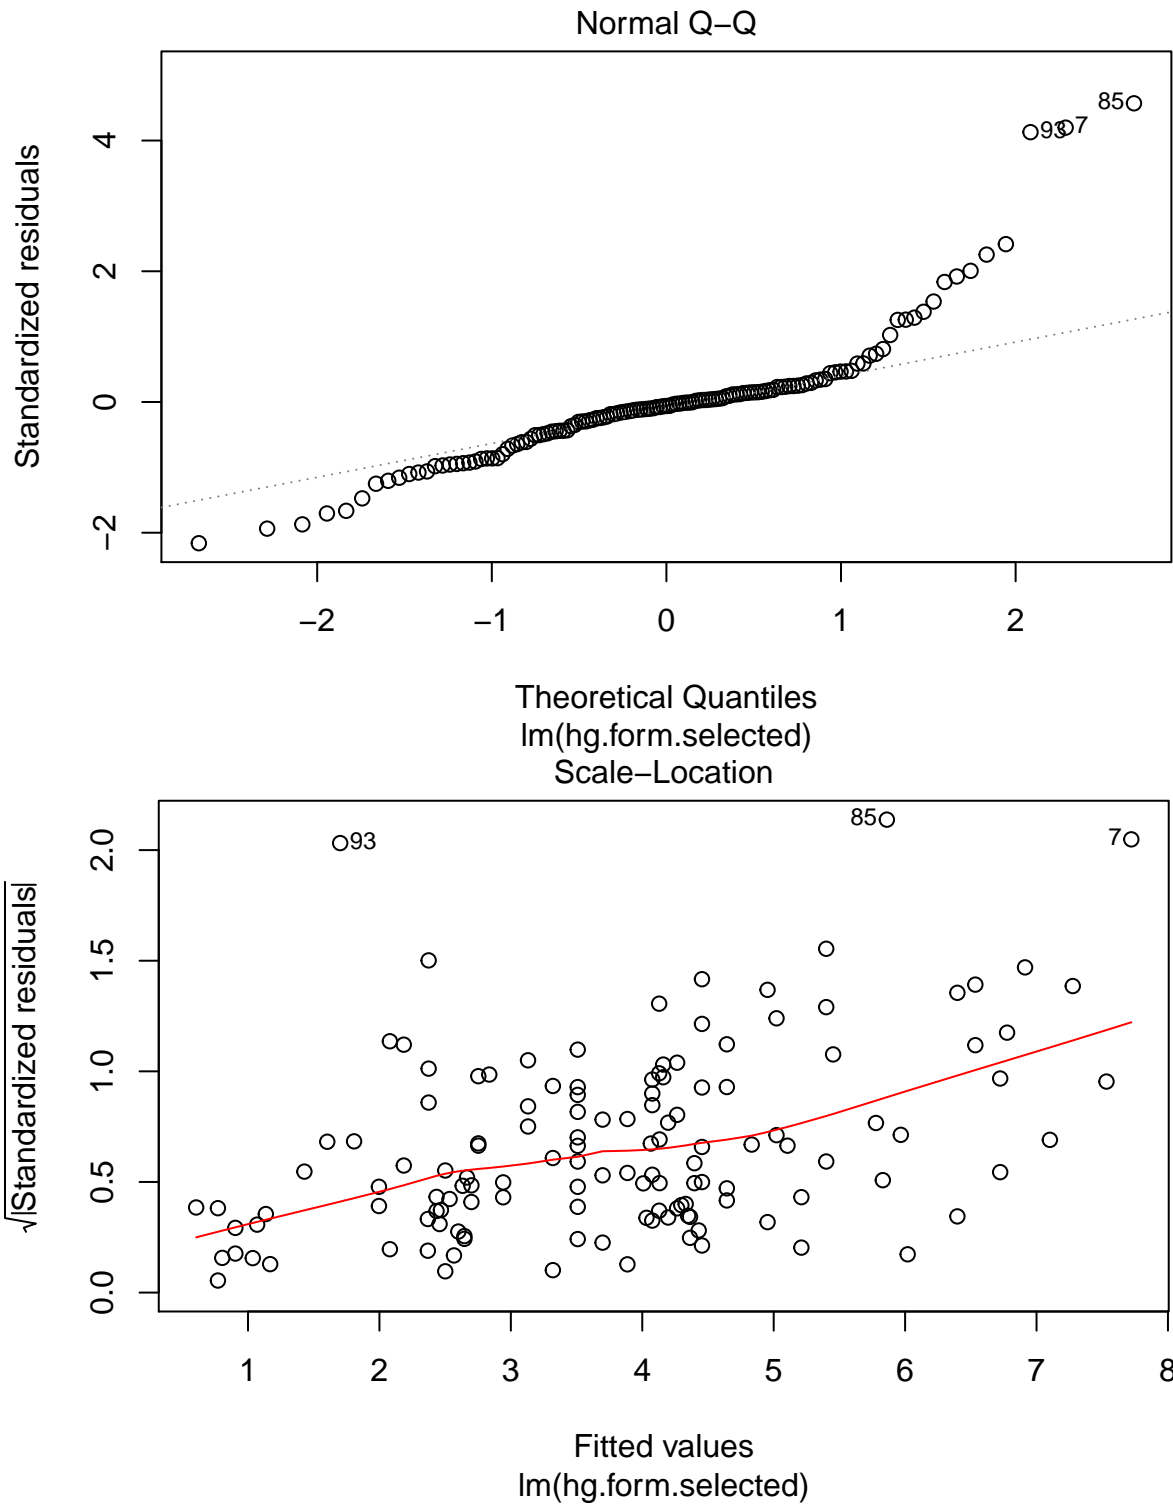


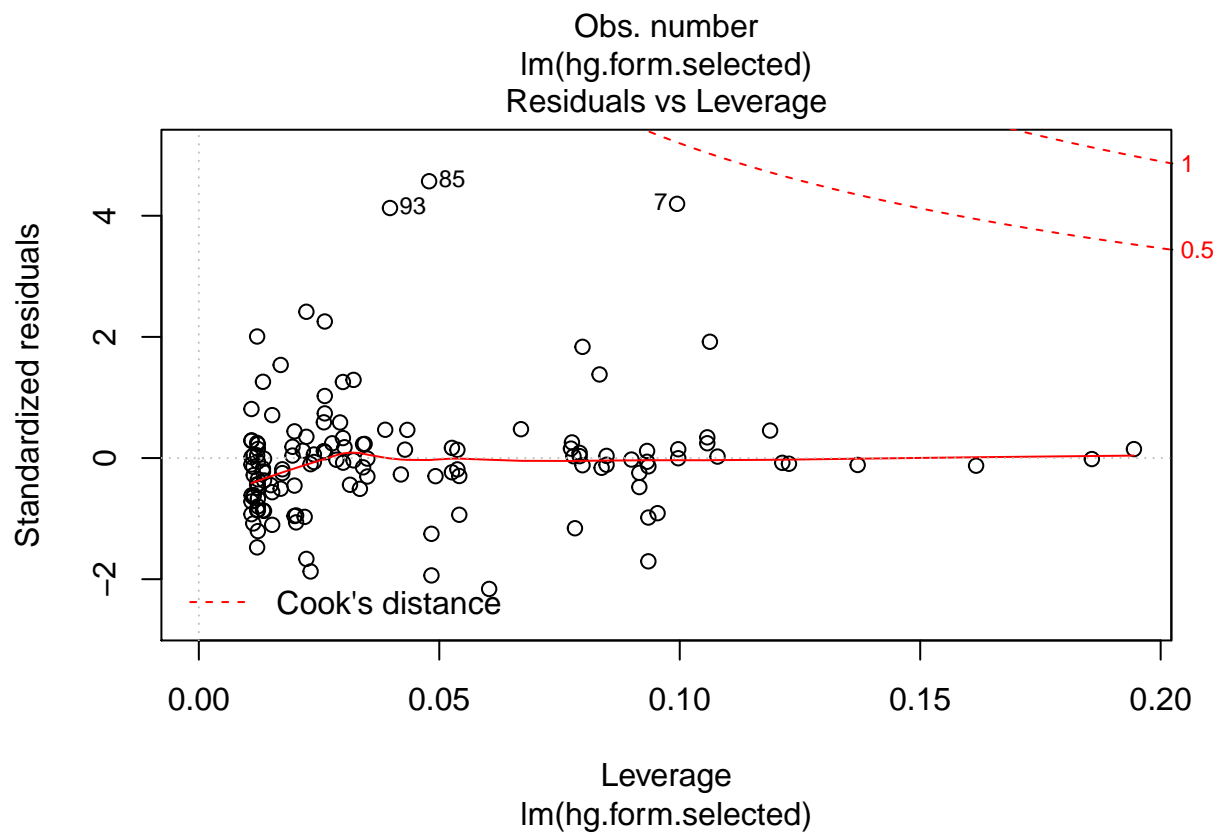

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.2935 -1.1455 -0.1474  0.5763 11.2698
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.41565     6.93854   -0.204   0.8387
## fisherman1      -8.99968     7.42906   -1.211   0.2280
## weight          0.03313     0.09907    0.334   0.7386
## fishmlwk        1.53107     0.70201    2.181   0.0310 *
## fisherman1:weight 0.15596     0.10556    1.477   0.1420
## fisherman1:fishmlwk -1.43279     0.70398   -2.035   0.0439 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.528 on 129 degrees of freedom
## Multiple R-squared:  0.2878, Adjusted R-squared:  0.2602
## F-statistic: 10.42 on 5 and 129 DF,  p-value: 2.042e-08
```


Commenter les résultats, valeurs des coefficients, font-elles du sens, p-values

Les diagnostic plots pour weight + dire que c'est pareil pour tout le monde









Conclusion (Lucile)

3 lignes

Les principales variables explicatives sont:

- La consommation de poisson => logique
- Le poids => plus étonnant mais on sait que les cellules graisseuses stockent les toxiques

References

Al-Majed, NB, and MR Preston. 2000. "Factors Influencing the Total Mercury and Methyl Mercury in the Hair of the Fishermen of Kuwait." *Environmental Pollution* 109 (2). Elsevier: 239–50.

Park, Jung-Duck, and Wei Zheng. 2012. "Human Exposure and Health Effects of Inorganic and Elemental Mercury." *Journal of Preventive Medicine and Public Health* 45 (6). Korean Society for Preventive Medicine: 344.